

1. Download entire folder from <http://tinyurl.com/tischR>
2. **IMPORTANT!** *Extract/unzip* all contents of folder to somewhere easy to find (desktop)
3. (optional) <http://collabedit.com/5fd9n>

Introduction to R

Josh Quan
Social Sciences Data Librarian
Tisch Library
Spring 2016

What we will cover in the hour:

R Studio Environment, Workflow, and Basics

Objects, Structures and Data types

Subsetting

R packages

Baby names dataset

Linear Modeling

Getting help

Why learn R?

- Second most commonly used scripting language, after SQL¹




Figure 1-10. Most commonly used tools

1. 2014 O'Reilly Data Science Salary Survey (<http://www.oreilly.com/data/free/files/2014-data-science-salary-survey.pdf>)

Why learn R?

- Second most commonly used scripting language, after SQL¹
 - Third most commonly used data tool after SQL and Excel¹
- R is popular in both industry and academia
- R is open source and has a large, active community
- Reproducible code = Reproducible workflow = Transparency
- In addition to advancing the needs of mathematicians/statisticians in ways that other tools can't, R is a bridge to “Data Science” competencies
 - Data manipulation, visualization, and machine learning

1. 2014 O'Reilly Data Science Salary Survey (<http://www.oreilly.com/data/free/files/2014-data-science-salary-survey.pdf>)

Back to earth...

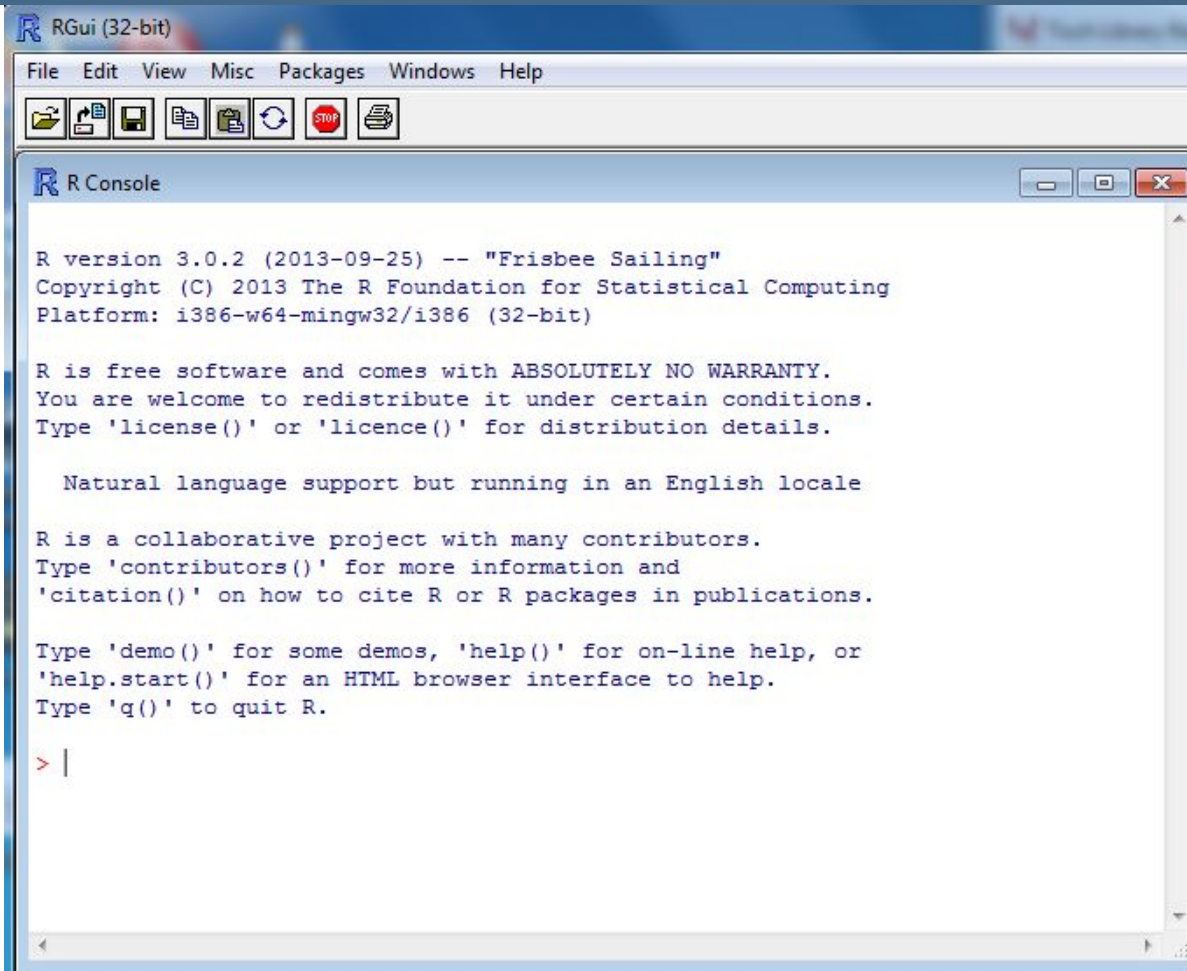
- R is not a data vault or spreadsheet. The easiest way to enter data in R is to enter it somewhere else, then import it. To visually inspect tabular data, SPSS or Excel are still better options.
- R makes some ordinary tasks more difficult right out of the box.
- The learning curve for R is nontrivial. Being a data analyst != being a programmer.
- Being open source is a boon and curse.
- R is not meant to be a “general” programming language like Python.

To be successful with R

- Like learning any programming language, take your time and try to run the code in your head before you run it on your machine. Try to predict what will happen.
- Be patient.
- Think of a fun project that you actually would like to do and do it in R.
- Ask your friendly librarian about useful resources.
- Understand that unlike excel, there are many paths to the same solution in R. You need not learn them all but to troubleshoot effectively and ask for help it is worthwhile to understand how others might work with R (ie., subsetting)
- If you master the techniques and concepts in this workshop you've mastered 80% of R. The rest is identifying specific packages/methodologies that are relevant to your [domain](#).

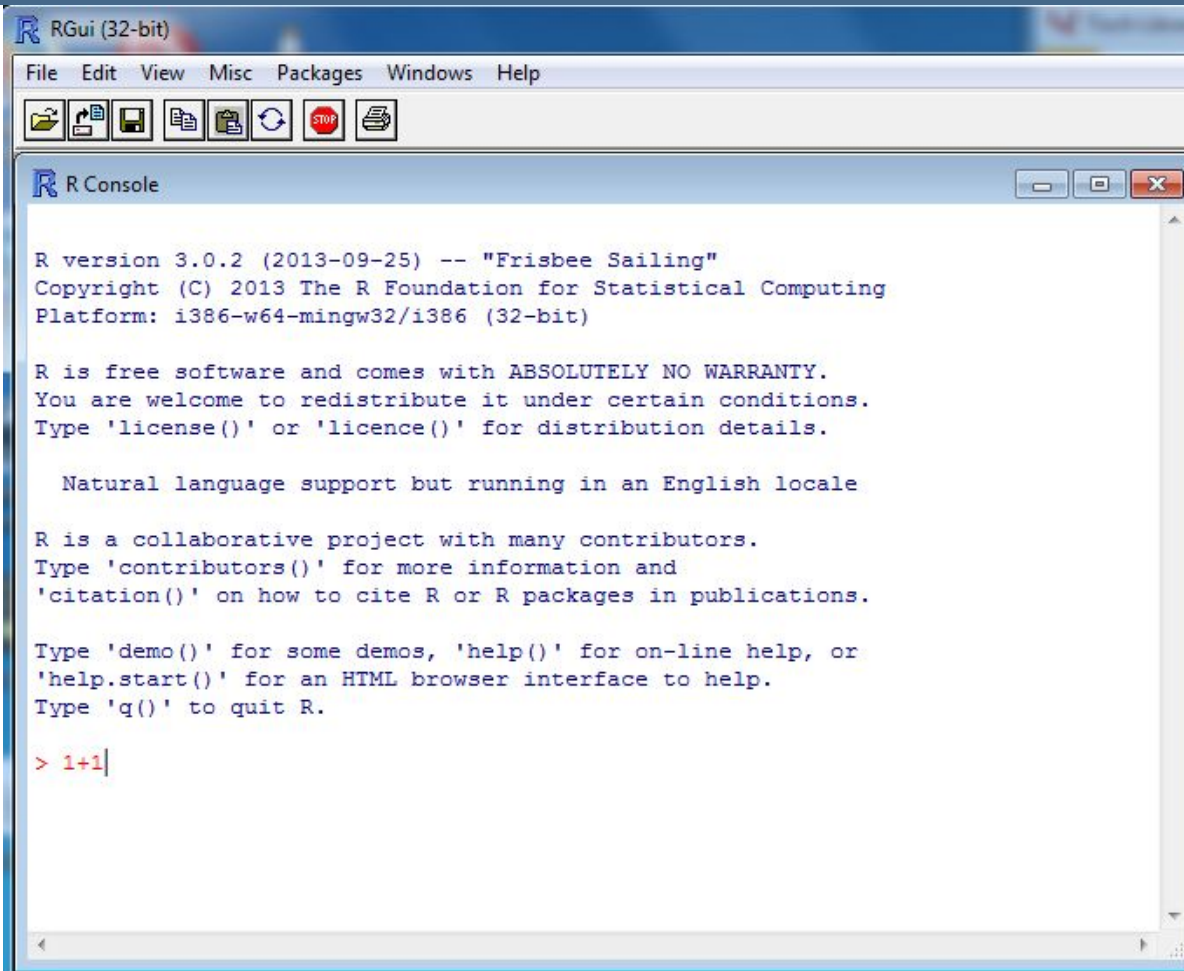
R Console

The console gives you a place to execute commands written in the R computer language.



R Console

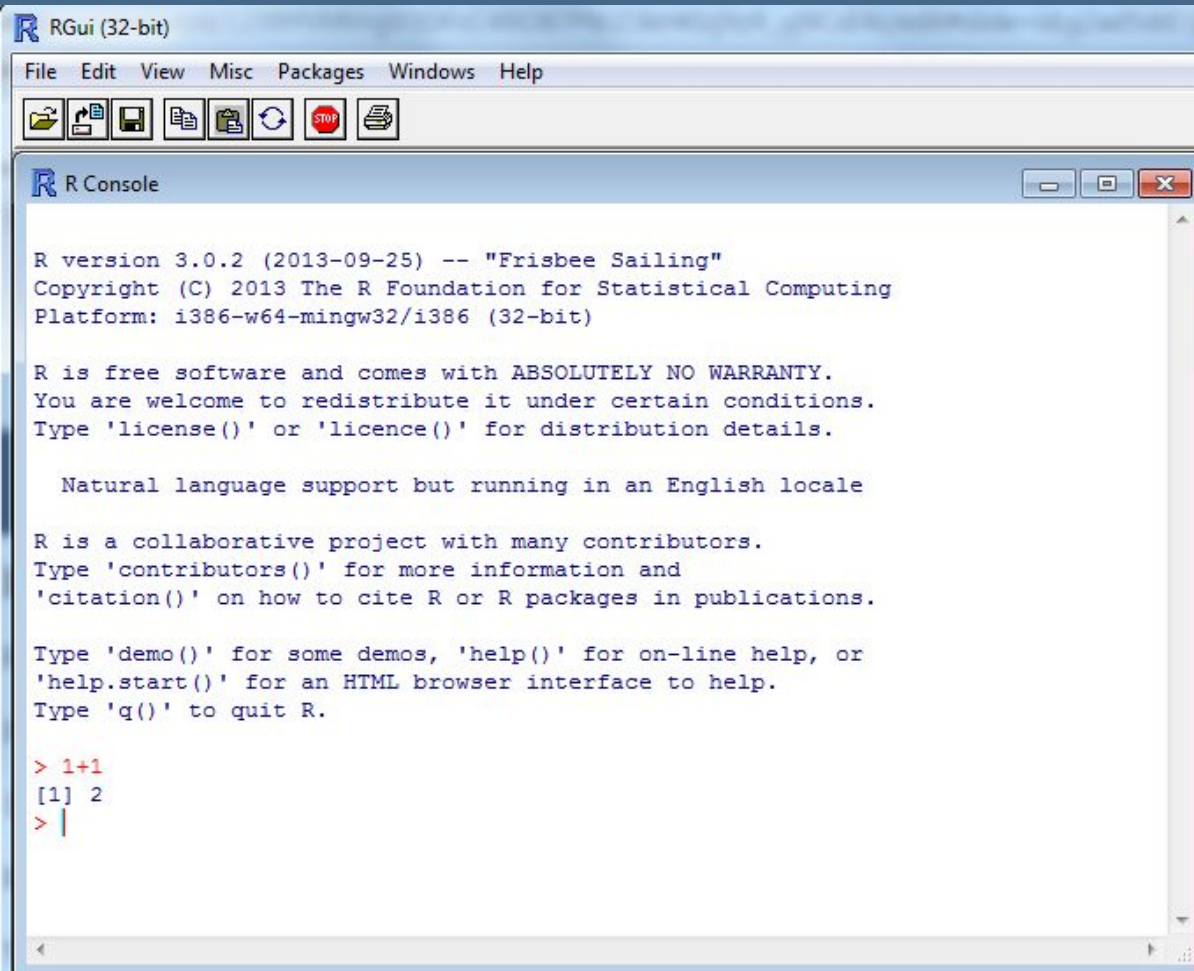
Type commands on the line that begins with a > sign (known as the prompt)



R Console

When you hit enter,
R will run your command
and display any output
below it

Output 
New Prompt 



The screenshot shows the RGui (32-bit) application window. The title bar reads "RGui (32-bit)". The menu bar includes "File", "Edit", "View", "Misc", "Packages", "Windows", and "Help". Below the menu bar is a toolbar with icons for file operations (open, save, print, etc.) and a "STOP" button. The main window is titled "R Console" and contains the following text:

```
R version 3.0.2 (2013-09-25) -- "Frisbee Sailing"
Copyright (C) 2013 The R Foundation for Statistical Computing
Platform: i386-w64-mingw32/i386 (32-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.

> 1+1
[1] 2
> |
```

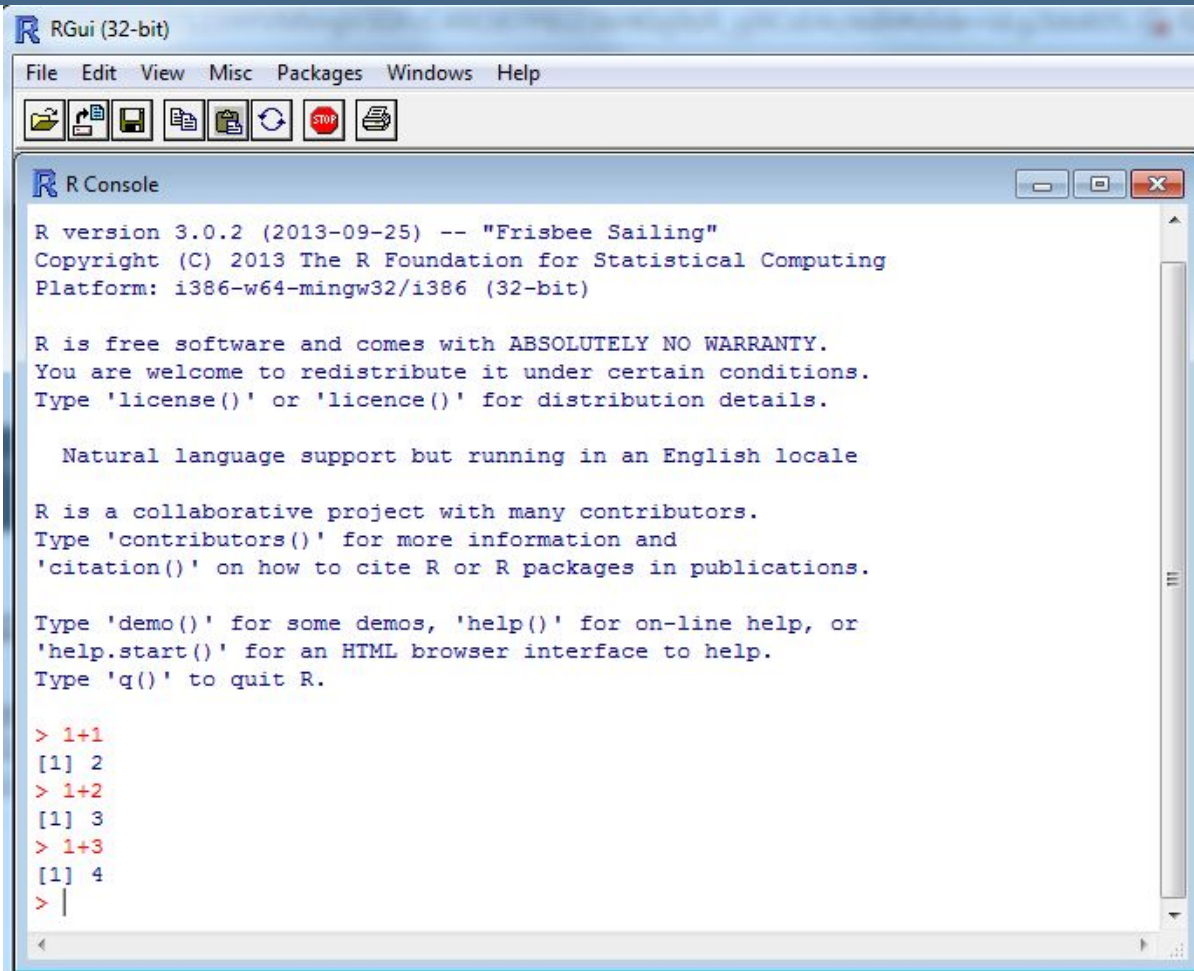
R Console

As you enter commands, you accumulate a history of past commands.

You can scroll through past commands by using the up arrow key on your keyboard.

R displays an index [1] next to the output.

Just ignore this.



The screenshot shows the RGui (32-bit) window. The title bar reads "RGui (32-bit)". The menu bar includes "File", "Edit", "View", "Misc", "Packages", "Windows", and "Help". Below the menu bar is a toolbar with icons for file operations and execution. The main window is titled "R Console" and contains the following text:

```
R version 3.0.2 (2013-09-25) -- "Frisbee Sailing"
Copyright (C) 2013 The R Foundation for Statistical Computing
Platform: i386-w64-mingw32/i386 (32-bit)

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

Natural language support but running in an English locale

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.
Type 'q()' to quit R.


> 1+1
[1] 2
> 1+2
[1] 3
> 1+3
[1] 4
> |
```

R is like a fancy calculator

$5 + 5$

$1 * 2$

$4 ^ 2$



R is like a fancy calculator

$5 + 5$


10

$1 * 2$

2

$4 ^ 2$

16




R can do algebra

```
a <- 1
```

```
b <- 2
```

```
a + b
```



R can do algebra

```
a <- 1
```


```
b <- 2
```

```
a + b
```

```
# 3
```

```
A <- 3
```

```
a + b - A
```



R can do algebra

```
a <- 1
```

```
b <- 2
```


```
a + b
```

```
# 3
```

```
A <- 3
```

```
a + b - A
```

```
# 0
```




It has built in functions that let you do more sophisticated manipulations

```
round(3.145)
```

```
factorial(3)
```

```
sqrt(9)
```

```
mean(c(7.5,8.2,3.1,5.6,10.9,4.6))
```



It has built in functions that let you do more sophisticated manipulations

```
round(3.145)
```

3

```
factorial(3)
```


6

```
sqrt(9)
```

3

```
mean(c(7.5,8.2,3.1,5.6,10.9,4.6))
```

6.65



Your Turn

What do you think this will return?

```
factorial(round(2.0015) + 1)
```

R always works from the innermost parenthesis to the outermost (just like a calculator)

`factorial(round(2.0015) + 1)`



`factorial (2 + 1)`



`factorial(3)`



6

+ prompt

if your prompt turns into a “+”,
R thinks you haven't finished your previous
command.

Either finish the command or press escape to
start over.

```
> 1+1
```

```
[1] 2
```

```
> 1+2
```

```
[1] 3
```

```
> 1+3
```

```
[1] 4
```

```
> factorial(round(2.0015)) + 1
```

```
+ |
```

+ prompt

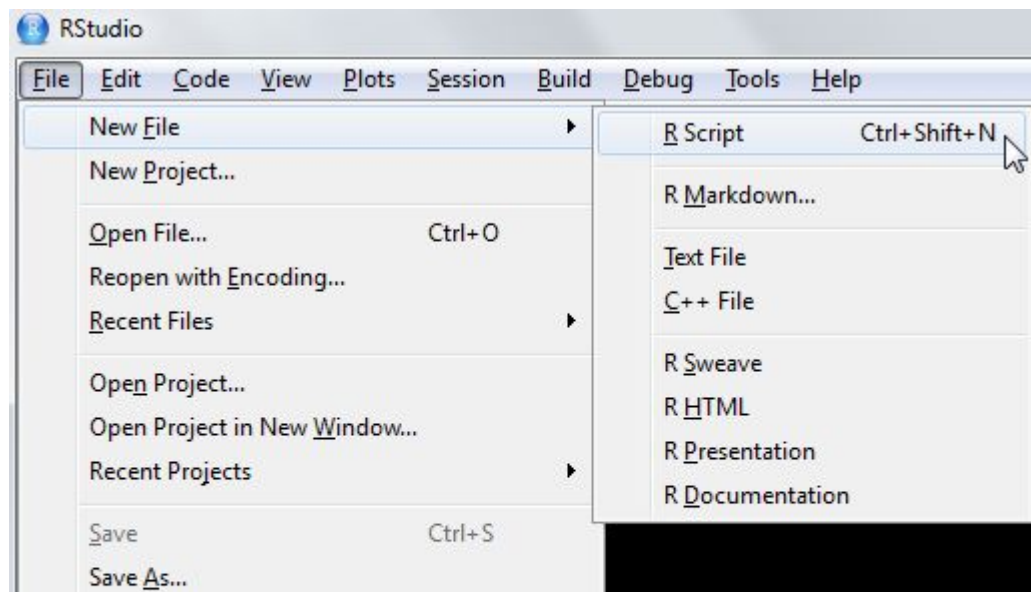
if your prompt turns into a “+”,
R thinks you haven’t finished your previous
command.

Either finish the command or press escape to
start over.

```
> 1+1
[1] 2
> 1+2
[1] 3
> 1+3
[1] 4
> factorial(round(2.0015) + 1
+ )
[1] 6
> |
```


R Scripts

It is much easier to compose your code in an R script than in the command line.
To open a script, go to File > New script in the toolbar



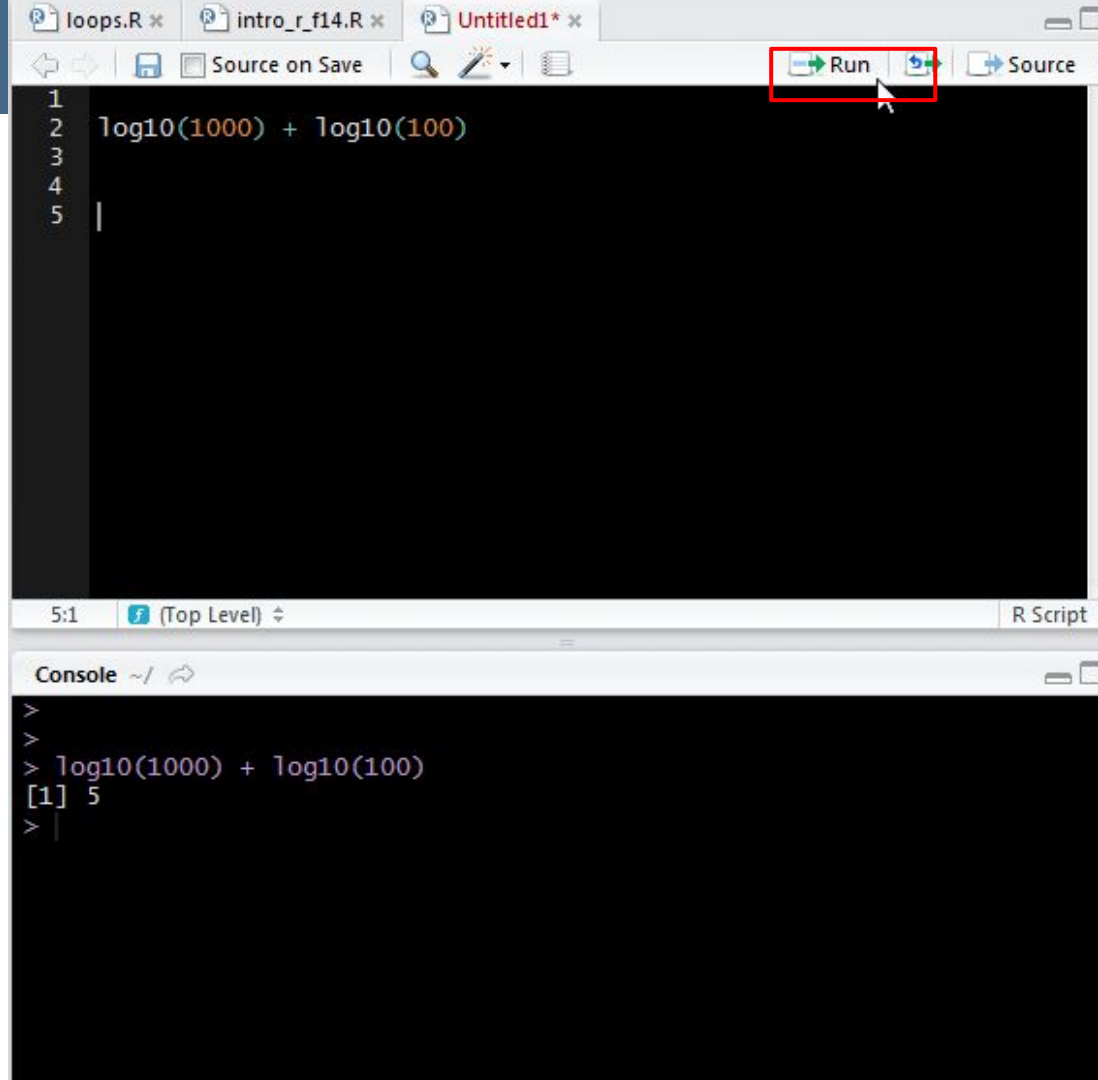
Common Workflow

1. Write code in an R script




Common Workflow

1. Write code in an R script
2. Run code in console with run icon or from script (Ctrl- r) or Ctrl-Enter



Common Workflow

1. Write code in an R script
2. Run code in console with run icon or from script (Ctrl- r) or Ctrl-Enter
3. Save R Script when finished



R Objects

+1 212 777 6640 №41710
\$129/£74/€108/¥15000 37**8
№10 ★1,150 ★2,275 ♦3,425
3*1.49+TAX \$4.85 #0632 22-46 TAX±8.625%
129.00ST 11.13TAX 140.13TL 19.87CG 491,305÷26

Save information as an R object with the less than sign followed by a minus, e.g, an arrow: <-

```
foo <- 42
```


Save information as an R object with the greater than sign followed by a minus, e.g, an arrow: <-



name of new object

foo <- 42

Save information as an R object with the greater than sign followed by a minus, e.g, an arrow: <-



**assignment
operator, “gets”**

foo <- 42

Save information as an R object with the greater than sign followed by a minus, e.g, an arrow: <-



**information to store
in the object**

foo <- 42

Common R Workflow

Save output of one function as an R object to use in a second function

```
foo <- round(3.1415) + 1
```

```
foo
```


Common R Workflow

Save output of one function as an R object to use in a second function

```
foo <- round(3.1415) + 1
```

```
foo
```

```
# 4
```

Common R Workflow

Save output of one function as an R object to use in a second function

```
foo <- round(3.1415) + 1
```

```
foo
```

```
# 4
```

```
factorial(foo)
```

Common R Workflow

Save output of one function as an R object to use in a second function

```
foo <- round(3.1415) +1
```

```
foo
```

```
# 4
```

```
factorial(foo)
```

```
# 24
```

rm

You can remove an object with rm

foo

4

```
rm(foo)
```

foo

rm

You can remove an object with rm

foo

4

```
rm(foo)
```

foo

Error: object 'foo' not found

This can be useful if you overwrite an object that comes with R

```
pi
```

```
# 3.141593
```

```
pi <- 1
```

```
pi
```

```
# 1
```

```
rm(pi)
```

```
pi
```

```
# 3.141593
```

Structures and Data Types

+1 212 777 6640 No 41710
\$129/£74/€108/¥15000 37**8
No 10 ★1,150 ★2,275 ♦3,425
3*1.49+TAX \$4.85 #0632 22-46 TAX+8.625%
129.00ST 11.13TAX 140.13TL 19.87CG 491,305÷26

Data Types

Like Excel, R can recognize different types of data

Four Basic Types:

Numeric

Character string (text)

Logical

Factor

Numeric

Any number, no quotes. Appropriate for math.

1 + 1

3000000

class(0.00001)

Numeric

Any number, no quotes. Appropriate for math.

1 + 1

3000000

class(0.00001)

"numeric"

Character

Any symbols surrounded by quotes.

Appropriate for words, variable names, messages, any text.

```
"hello"
```

```
class("hello")
```

Character

Any symbols surrounded by quotes.

Appropriate for words, variable names, messages, any text.

“hello”

class(“hello”)

"character"

“hello” + “world”

Error

nchar("hello")

5

paste("hello", "world")

“hello world”

Logical

TRUE or FALSE. R's form of binary data. Useful for logical tests.

```
3 > 4
```

```
# FALSE
```

```
class(TRUE)
```

```
# "logical"
```

```
class (T)
```

```
# "logical"
```

Factor

R's form of categorical data. Saved as an integer with a set of labels (e.g. levels)

```
fac <- factor(c("a", "b", "c"))
```

```
fac
```

```
# a b c
```

```
# Levels: abc
```

```
class(fac)
```

```
# factor
```

Type	Examples
numeric	0, 1, -2, 3.14, 0.0005
character	"gender", "date", "31"
logical	TRUE, FALSE, T, F
factor	a c c b Levels: abc

Structures

Five types of structures in R:

Vectors

Matrices

Arrays

Lists

Data frames

Structures

Five types of structures in R:

Vectors

Matrices

Arrays

Lists

Data frames

Vectors

Combine multiple elements into a one dimensional array

Create with the c function

```
vec <- c(1, 2, 3, 10, 100)
```

Vectors

Combine multiple elements into a one dimensional array

Create with the c function

```
vec <- c(1, 2, 3, 10, 100)  
vec
```

```
# 1 2 3 10 100
```



Matrices

Multiple elements stored in a two dimensional array

Create with the matrix function

```
mat <- matrix(c(1, 2, 3, 4, 5, 6), nrow = 2)
```

Matrices

Multiple elements stored in a two dimensional array

Create with the matrix function

```
mat <- matrix(c(1, 2, 3, 4, 5, 6), nrow = 2)  
mat
```

```
#   [,1] [,2] [,3]  
# [1,]  1  3  5  
# [2,]  2  4  6
```

Matrices

Multiple elements stored in a two dimensional array

Create with the matrix function

```
mat <- matrix(c(1, 2, 3, 4, 5, 6), nrow = 2)  
mat
```

```
#   [,1] [,2] [,3]  
# [1,]  1  3  5  
# [2,]  2  4  6
```



**vector of elements
to go in the matrix**

```
mat <- matrix(c(1, 2, 3, 4, 5, 6), nrow = 2)  
mat
```

```
#   [,1] [,2] [,3]  
# [1,]  1  3  5  
# [2,]  2  4  6
```


number of rows for
matrix

```
mat <- matrix(c(1, 2, 3, 4, 5, 6), nrow = 2)
```

```
mat
```

```
#      [,1] [,2] [,3]  
# [1,]  1   3   5  
# [2,]  2   4   6
```

```
mat <- matrix(c(1, 2, 3, 4, 5, 6), nrow = 3)
```

```
mat <- matrix(c(1, 2, 3, 4, 5, 6), nrow = 3)  
mat
```

```
      [,1] [,2]  
[1,]    1    4  
[2,]    2    5  
[3,]    3    6
```

```
mat <- matrix(c(1, 2, 3, 4, 5, 6), nrow = 3, byrow = TRUE)  
mat
```

```
      [,1] [,2]  
[1,]    1    2  
[2,]    3    4  
[3,]    5    6
```

```
vec <- c(1, 2, 3, "10", 100)
```

```
vec <- c(1, 2, 3, "10", 100)
```

```
class(vec)
```

```
# "character"
```

Vectors and Matrices only allow one data type.

So what do we do if we want to work with multiple data types?

Data frame

A data frame is a two dimensional group of R objects.

Each column in a data frame can be a different type

```
df <- data.frame(c(1, 2, 3),  
c("R","S","T"), c(TRUE, FALSE, TRUE))
```

```
class (df)
```


Data frame

A data frame is a two dimensional group of R objects.

Each column in a data frame can be a different type

```
df <- data.frame(c(1, 2, 3),  
c("R","S","T"), c(TRUE, FALSE, TRUE))
```

```
class(df)
```

```
# "data.frame"
```

names

You can name the elements of a vector, list or data frame when you create them.

```
nvec <- c(one = 1, two = 2, three = 3)
```

```
nvec
```

```
# one two three
```

```
# 1 2 3
```

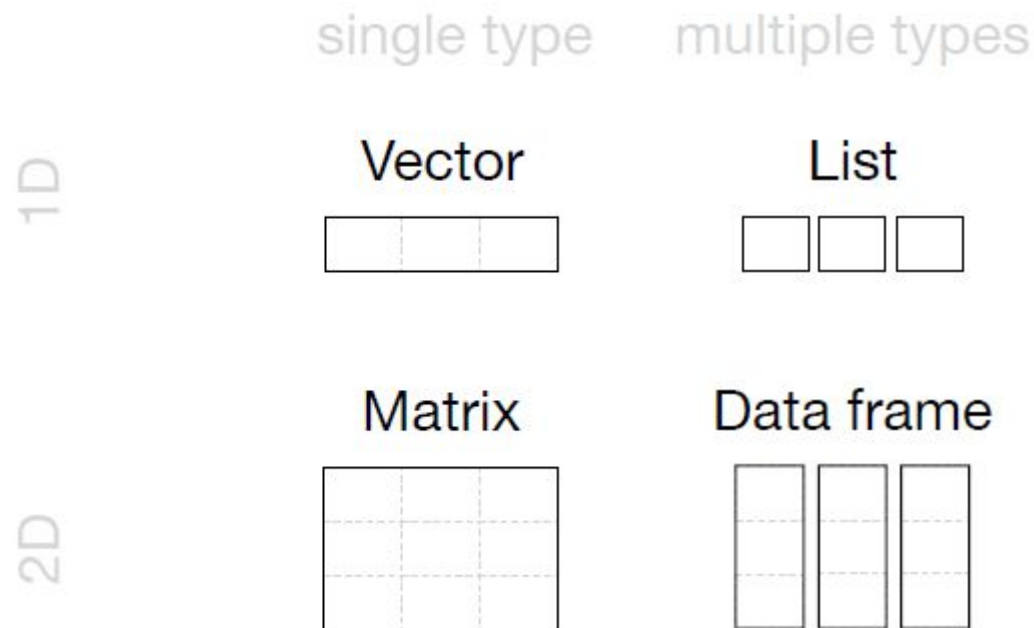
```
ndf <- data.frame(numbers = c(1, 2, 3),  
letters = c("R","S","T"),  
logic = c(TRUE, FALSE, TRUE))
```

```
ndf
```

```
ndf <- data.frame(numbers = c(1, 2, 3),  
letters = c("R","S","T"),  
logic = c(TRUE, FALSE, TRUE))
```

ndf

```
# numbers letters logic  
# 1 1 R TRUE  
# 2 2 S FALSE  
# 3 3 T TRUE
```



```
x <- c(0, 0, 0, 0, 1, 0 ,0)
y <- x
y
# 0 0 0 0 1 0 0
```

How can you save just the fifth element of x to y?

How can you change the fifth element of x to a 0?

+1 212 777 6640 №41710

Subsetting

\$129/£74/€108/¥15000 37***8

№10 ★1,150 ★2,275 ♦3,425

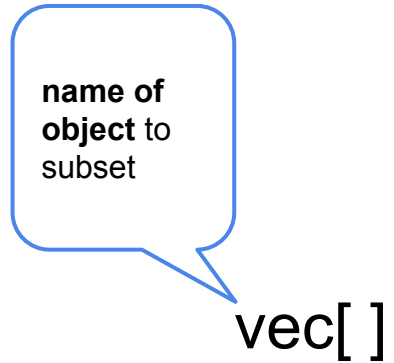
3*1.49+TAX \$4.85 #0632 22-46 TAX±8.625%

129.00ST 11.13TAX 140.13TL 19.87CG 491,305÷26

Subset notation

`vec[]`

Subset notation




A diagram illustrating the subset notation. A blue speech bubble with a tail pointing to the text "vec[]" contains the text "name of object to subset".

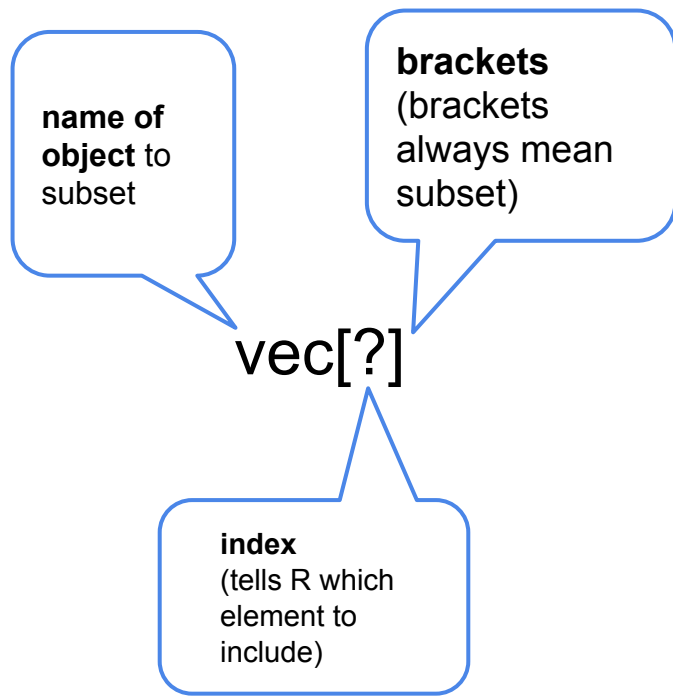
name of
object to
subset

vec[]

Subset notation



Subset notation



Each dimension needs its own index

vec [?]

6	1	3	6	10	5
---	---	---	---	----	---

Each dimension needs its own index

vec [?]

6	1	3	6	10	5
---	---	---	---	----	---

Each dimension needs its own index

df [?,?]

John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Each dimension needs its own index

df [?,?]

which **rows**
to include

John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Each dimension needs its own index

df [?,?]

which **rows**
to include

which **columns**
to include

John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Each dimension needs its own index

df [?,?]

which **rows**
to include

which **columns**
to include

separate
with comma

John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

```
vec <- c(6, 1, 3, 6, 10, 5)
```

```
df <- data.frame(  
  name = c("John", "Paul", "George", "Ringo"),  
  birth = c(1940, 1942, 1943, 1940),  
  instrument = c("guitar", "bass", "guitar", "drums")  
)
```

run the code on the following slide **IN YOUR HEADS**

vec

6	1	3	6	10	5
---	---	---	---	----	---

df

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Predict what the following code will do

DON'T RUN IT

```
vec[2]  
vec[c(5, 6)]  
vec[-c(5,6)]  
vec[vec > 5]
```

```
df[c(2, 4), 3]  
df[ , 1]  
df[ , "instrument"]  
df$instrument
```

Four ways to subset using []

1. Integers
2. Blank spaces
3. Names
4. Logical

Integers (positive)

positive integers behave just like *ij* notation in linear algebra


df[?,?]

John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Integers (positive)

positive integers behave just like *ij* notation in linear algebra

df[2,?]

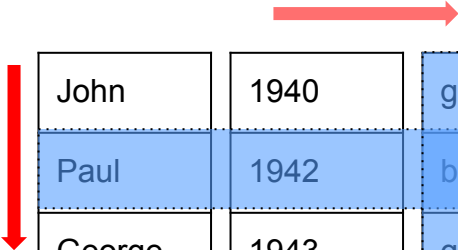


John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Integers (positive)

positive integers behave just like *ij* notation in linear algebra

df[2,3]



John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Integers (positive)

positive integers behave just like *ij* notation in linear algebra

df[2,3]

John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Integers (positive)

positive integers behave just like *ij* notation in linear algebra

df[c(2,4), ?]

John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Integers (positive)

positive integers behave just like *ij* notation in linear algebra

`df[c(2,4), c(2,3)]`

John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Integers (positive)

positive integers behave just like *ij* notation in linear algebra

`df[c(2,4), c(2,3)]`

John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Integers (positive)

positive integers behave just like *ij* notation in linear algebra

`df[c(2,4), 3]`

John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

?

Integers (positive)

positive integers behave just like *ij* notation in linear algebra

`df[c(2,4), 3]`

John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Integers (positive)

Colons are a useful way to create vectors and to return results

```
df[1:4, 1:2]
```

John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Integers (negative)

Negative integers return **everything but** the elements at the specified locations.

You cannot use both negative and positive integers in the same direction

`vec[-c(5,6)]`

6	1	3	6	10	5
---	---	---	---	----	---

Integers (negative)

Negative integers return **everything but** the elements at the specified locations.

You cannot use both negative and positive integers in the same direction

`vec[-c(5,6)]`

6	1	3	6	10	5
---	---	---	---	----	---

Blank spaces

Blank spaces return **everything**

(i.e., no subsetting occurs on that dimension)

vec[]

6	1	3	6	10	5
---	---	---	---	----	---

Blank spaces

Blank spaces return **everything**

(i.e., no subsetting occurs on that dimension)

vec[]

6	1	3	6	10	5
---	---	---	---	----	---

Blank spaces

Blank spaces return **everything**

(i.e., no subsetting occurs on that dimension)

`df[1,]`

John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Blank spaces

Blank spaces return **everything**

(i.e., no subsetting occurs on that dimension)

`df[,2]`

John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Names

If your object has names, you can ask for elements or columns back by name.

`vec[]`

6	1	3	6	10	5
---	---	---	---	----	---

Names

If your object has names, you can ask for elements or columns back by name.

```
names(vec) <- c("a","b","c","d","e","f")
```

```
vec[  ]
```

a	b	c	d	e	f
6	1	3	6	10	5

Names

If your object has names, you can ask for elements or columns back by name.

```
names(vec) <- c("a","b","c","d","e","f")
```

```
vec[c("a","b","d")]
```

a	b	c	d	e	f
6	1	3	6	10	5

Names

If your object has names, you can ask for elements or columns back by name.

```
df[, "birth"]
```

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Names

If your object has names, you can ask for elements or columns back by name.

```
df[,c("name","birth")]
```

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Names \$

A common syntax for subsetting lists and data frames with names is the dollar sign.

`df$birth`

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Names \$

A common syntax for subsetting lists and data frames with names is the dollar sign.

df\$birth

name of
dataframe

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Names \$

A common syntax for subsetting lists and data frames with names is the dollar sign.

df\$birth

\$

name of
dataframe

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Names \$

A common syntax for subsetting lists and data frames with names is the dollar sign.

df\$birth

\$

name of
dataframe

name of
column
(no quotes)

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Names \$

A common syntax for subsetting lists and data frames with names is the dollar sign.

`df$birth`

\$

name of
dataframe

name of
column
(no quotes)

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Logical

You can subset with a logical vector of the same length as the dimension you are subsetting. Each element that corresponds to a **TRUE** will be returned.

```
vec[c(FALSE,TRUE,FALSE,TRUE,TRUE,FALSE)]
```

6	1	3	6	10	5
---	---	---	---	----	---

Logical

You can subset with a logical vector of the same length as the dimension you are subsetting. Each element that corresponds to a **TRUE** will be returned.

```
vec[c(FALSE,TRUE,FALSE,TRUE,TRUE,FALSE)]
```

6	1	3	6	10	5
---	---	---	---	----	---

Logical operators

operator	tests
<code>x > y</code>	is x greater than y?
<code>x >= y</code>	is x greater than or equal to y?
<code>x < y</code>	is x less than y?
<code>x <= y</code>	is x less than or equal to y?
<code>x == y</code>	is x equal to y?
<code>x != y</code>	is x not equal to y?
<code>x %in% c(y,z)</code>	is x in the set c(y, z)?

Boolean operators

operator	tests
a & b	both a and b are TRUE
a b	at least one of a and b is TRUE (or)
xor(a,b)	a is TRUE or b is TRUE, but not both
!(a)	not a (TRUE goes to FALSE, FALSE goes to TRUE)
any(a,b,c)	at least one of a, b , or c is TRUE
all(a,b,c)	each of a, b, and c is TRUE

Logical tests(standard and boolean operators)

Combining logical tests with subsetting is a very powerful technique!

```
df[df$instrument == "guitar", ]
```

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Logical tests(standard and boolean operators)

Combining logical tests with subsetting is a very powerful technique!

```
df[df$instrument == "guitar", ]
```

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Logical tests(standard and boolean operators)

Combining logical tests with subsetting is a very powerful technique!

```
df[df$birth > 1940, 1]
```

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Logical tests(standard and boolean operators)

Combining logical tests with subsetting is a very powerful technique!

```
df[df$birth > 1940, 1]
```

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Logical tests(standard and boolean operators)

Combining logical tests with subsetting is a very powerful technique!

```
df[df$birth < 1943 & df$instrument != "guitar", ]
```

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums


Logical tests(standard and boolean operators)

Combining logical tests with subsetting is a very powerful technique!

```
df[df$birth < 1943 & df$instrument != "guitar", ]
```

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

BUT WAIT



THERE'S MORE

Subset() Function

The easiest way to subset is to use the `subset()` function. Notice we don't use brackets `[]`. The form of the subset function is,

`subset(x, subset, select)`

`subset(df, birth > 1940, select = name)`

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Subset() Function

The easiest way to subset is to use the `subset()` function. Notice we don't use brackets `[]`. The form of the subset function is,

`subset(x, subset, select)`

`subset(df, birth > 1940, select = name)`

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Subset() Function

The easiest way to subset is to use the `subset()` function. Notice we don't use brackets `[]`. The form of the subset function is,

`subset(x, subset, select)`

`subset(df, name == "Paul")`

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Subset() Function

The easiest way to subset is to use the `subset()` function. Notice we don't use brackets `[]`. The form of the subset function is,

`subset(x, subset, select)`

`subset(df, name == "Paul")`

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Subset() Function

The easiest way to subset is to use the `subset()` function. Notice we don't use brackets `[]`. The form of the subset function is,

`subset(x, subset, select)`

`subset(df, select = birth)`

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

Subset() Function

The easiest way to subset is to use the `subset()` function. Notice we don't use brackets `[]`. The form of the subset function is,

`subset(x, subset, select)`

`subset(df, select = birth)`

name	birth	instrument
John	1940	guitar
Paul	1942	bass
George	1943	guitar
Ringo	1940	drums

	effect
integers	positive: returns specified elements ----- negative: returns everything but the specified elements
blank spaces	returns everything
names	returns elements or columns with the specified names
logicals	returns elements that correspond to TRUE. Can use logical tests with booleans/standard operators
subset()	returns what is specified in subset and/or select arguments

R packages

+1 212 777 6640 №41710
\$129/£74/€108/¥15000 37**8
№10 ★1,150 ★2,275 ♦3,425
3*1.49+TAX \$4.85 #0632 22-46 TAX+8.625%
129.00ST 11.13TAX 140.13TL 19.87CG 491,305÷26

R Packages

A collection of functions written for the R language.

Usually focuses on a specific task or problem.


Most of the useful R applications appear in packages.

There are over 5100 R packages

Top 10 R packages from Jan-May 2013

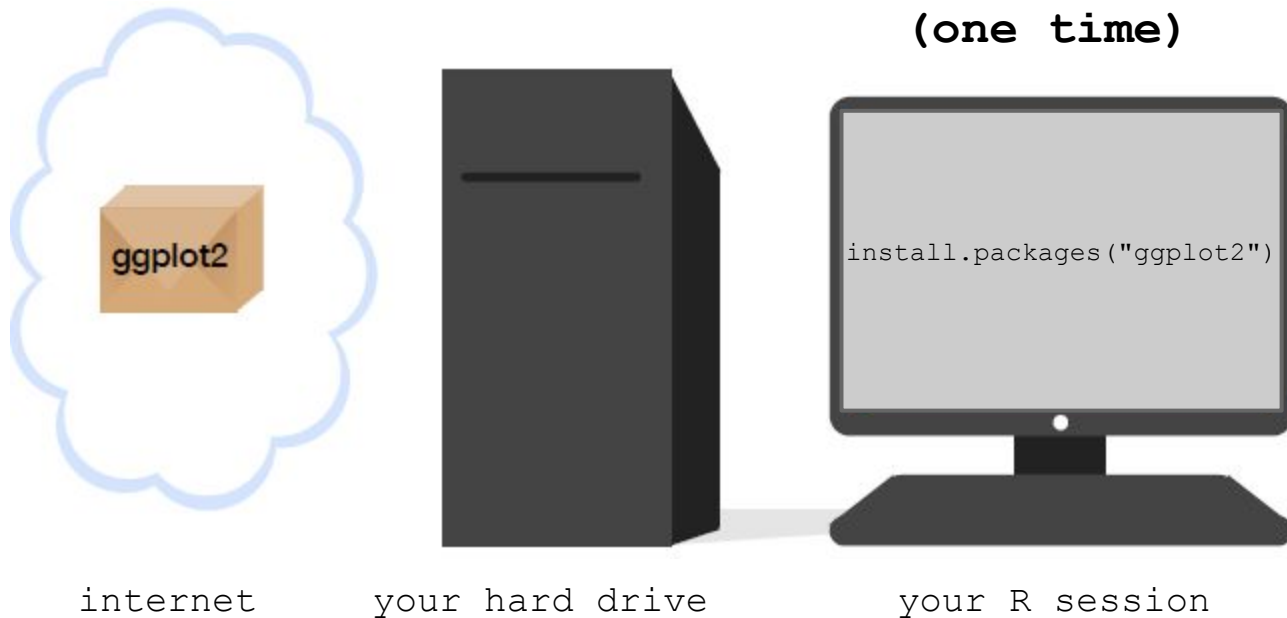
source: <http://www.r-statistics.com/>

	PACKAGE	TITLE	DOWNLOADS
1	plyr	Tools for splitting, applying and combining data	84049
2	digest	Create cryptographic hash digests of R objects	83192
3	ggplot2	An implementation of the Grammar of Graphics	82768
4	colorspace	Color Space Manipulation	81901
5	stringr	Make it easier to work with strings	77658
6	RColorBrewer	ColorBrewer palettes	66783
7	reshape2	Flexibly reshape data: a reboot of the reshape package	64911
8	zoo	S3 Infrastructure for Regular and Irregular Time Series (Z's ordered observations)	60844
9	proto	Prototype object-based programming	59043
10	scales	Scale functions for graphics	58369



Packages

```
install your package with  
install.packages("ggplot2")
```



Load your package with
`library(ggplot2)`





Baby names dataset

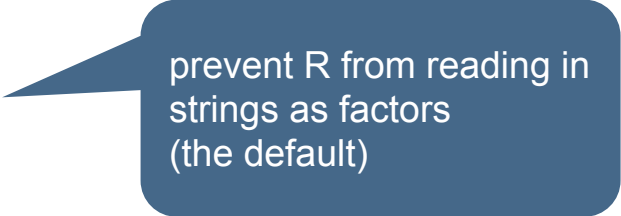
**Top 1000 female and male baby names in the US,
from 1880 to 2008.**

258,000 records ($1000 * 2 * 129$)

Five variables: year, name, soundex, sex, prop


```
library(plyr)  
library(ggplot2)
```

```
options(stringsAsFactors = FALSE)
```



prevent R from reading in
strings as factors
(the default)

```
bnames <- read.csv("bnames.csv")
```



make sure the data sets are in
your working directory


```
library(plyr)
```

Function	Package
subset	base
mutate	plyr
arrange	plyr

They all share similar syntax. The first argument is a data frame, and all other arguments are interpreted in the context of that data frame. Each returns a data frame.

And yes...another way to subset.

head(bnames)

	X	v1	v2	v3	v4	v5
1	1	1880	John	0.081541	boy	J500
2	2	1880	William	0.080511	boy	W450
3	3	1880	James	0.050057	boy	J520
4	4	1880	Charles	0.045167	boy	C642
5	5	1880	George	0.043292	boy	G620
6	6	1880	Frank	0.027380	boy	F652

tail(bnames)

	X	v1	v2	v3	v4	v5
257995	257995	2008	Diya	0.000128	girl	D000
257996	257996	2008	Carleigh	0.000128	girl	C642
257997	257997	2008	Iyana	0.000128	girl	I500
257998	257998	2008	Kenley	0.000127	girl	K540
257999	257999	2008	Sloane	0.000127	girl	S450
258000	258000	2008	Elianna	0.000127	girl	E450

View(bnames)

fix(bnames)

	X	v1	v2	v3	v4	v5
1	1	1880	John	0.081541	boy	J500
2	2	1880	William	0.080511	boy	W450
3	3	1880	James	0.050057	boy	J520
4	4	1880	Charles	0.045167	boy	C642
5	5	1880	George	0.043292	boy	G620
6	6	1880	Frank	0.02738	boy	F652
7	7	1880	Joseph	0.022229	boy	J210
8	8	1880	Thomas	0.021401	boy	T520
9	9	1880	Henry	0.020641	boy	H560
10	10	1880	Robert	0.020404	boy	R163
11	11	1880	Edward	0.019965	boy	E363
12	12	1880	Harry	0.018175	boy	H600
13	13	1880	Walter	0.014822	boy	W436
14	14	1880	Arthur	0.013504	boy	A636
15	15	1880	Fred	0.013251	boy	F630
16	16	1880	Albert	0.012609	boy	A416
17	17	1880	Samuel	0.008648	boy	S540
18	18	1880	David	0.007339	boy	D130
19	19	1880	Louis	0.006993	boy	L200

```
bnames$X <- NULL #drop extraneous X variable  
head(bnames)
```

	v1	v2	v3	v4	v5
1	1880	John	0.081541	boy	J500
2	1880	William	0.080511	boy	W450
3	1880	James	0.050057	boy	J520
4	1880	Charles	0.045167	boy	C642
5	1880	George	0.043292	boy	G620
6	1880	Frank	0.027380	boy	F652

#rename variables. limitation here is that you have to enter all of them in order. to rename just a single variable, the rename function is faster (reshape package)

```
names(bnames) <- c("year", "name", "prop", "sex", "soundex")  
head(bnames)
```

	year	name	prop	sex	soundex
1	1880	John	0.081541	boy	J500
2	1880	William	0.080511	boy	W450
3	1880	James	0.050057	boy	J520
4	1880	Charles	0.045167	boy	C642
5	1880	George	0.043292	boy	G620
6	1880	Frank	0.027380	boy	F652

summary(bnames)

year	name	prop	sex
Min. :1880	Jessie : 258	Min. :0.0000260	boy :129000
1st Qu.:1912	Leslie : 247	1st Qu.:0.0000810	girl:129000
Median :1944	Guadalupe: 244	Median :0.0001640	
Mean :1944	Jean : 244	Mean :0.0008945	
3rd Qu.:1976	Lee : 240	3rd Qu.:0.0005070	
Max. :2008	James : 239	Max. :0.0815410	
	(Other) :256528		

soundex
J500 : 4693
D500 : 3177
L500 : 2445
L200 : 2288
J200 : 1953
R200 : 1893
(Other):241551

```
joshua <- subset(bnames, name == "Joshua")  
head(joshua)
```

	year	name	prop	sex	soundex
212	1880	Joshua	0.000481	boy	J200
1249	1881	Joshua	0.000369	boy	J200
2236	1882	Joshua	0.000410	boy	J200
3202	1883	Joshua	0.000489	boy	J200
4266	1884	Joshua	0.000334	boy	J200
5254	1885	Joshua	0.000371	boy	J200

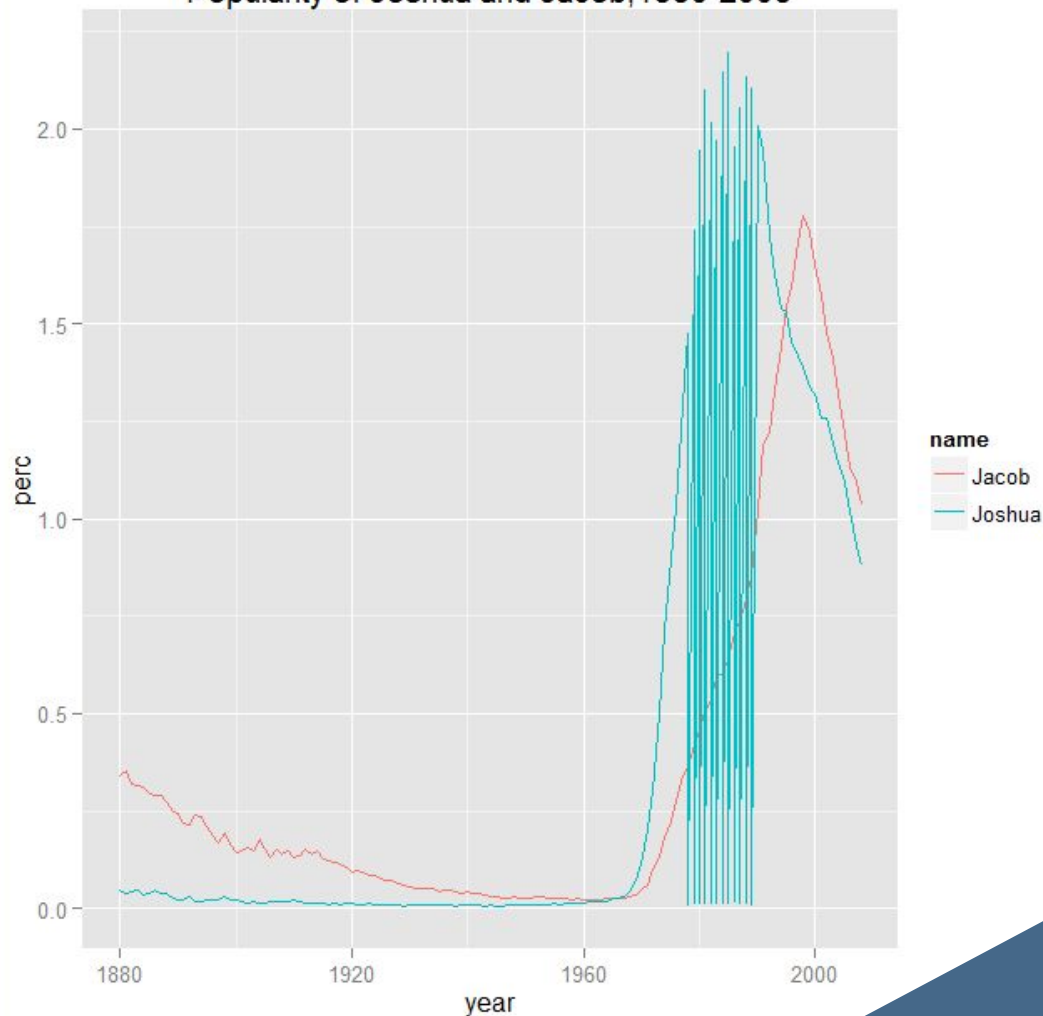

```
joshua <- mutate(joshua, perc = prop * 100)    #create new variable  
head(joshua)
```

	year	name	prop	sex	soundex	perc
1	1985	Joshua	0.021946	boy	J200	2.1946
2	1984	Joshua	0.021468	boy	J200	2.1468
3	1988	Joshua	0.021324	boy	J200	2.1324
4	1989	Joshua	0.021047	boy	J200	2.1047
5	1981	Joshua	0.020980	boy	J200	2.0980
6	1987	Joshua	0.020525	boy	J200	2.0525

```
joshua <- arrange(joshua, desc(perc))  
head(joshua, n=10)
```


	year	name	prop	sex	soundex	perc
1	1985	Joshua	0.021946	boy	J200	2.1946
2	1984	Joshua	0.021468	boy	J200	2.1468
3	1988	Joshua	0.021324	boy	J200	2.1324
4	1989	Joshua	0.021047	boy	J200	2.1047
5	1981	Joshua	0.020980	boy	J200	2.0980
6	1987	Joshua	0.020525	boy	J200	2.0525
7	1982	Joshua	0.020160	boy	J200	2.0160
8	1990	Joshua	0.020095	boy	J200	2.0095
9	1983	Joshua	0.019708	boy	J200	1.9708
10	1986	Joshua	0.019550	boy	J200	1.9550

Popularity of Joshua and Jacob, 1880-2008




ggplot2

Popularity of Joshua and Jacob, 1880-2008




```
joshua <- subset(bnames, name == "Joshua") #subset Joshua
```

```
qplot(year, perc, data = joshua, geom = "line") #quick plot
```




```
qplot(year, perc, data = joshua, geom = "point")  #quick plot
```




```
qplot(year, perc, data = joshua, geom = "point", color = sex)
```


creates a
different
colored set of
points for
each group of
sex
(male, female)




```
joshua <- subset(joshua, name == "Joshua" & sex == "boy")
```

```
qplot(year, perc, data = joshua, geom = "line") + ggtitle("Popularity of Joshua, 1880-2008")
```

Popularity of Joshua, 1880-2008



Popularity of Joshua and Jacob, 1880-2008



Your Turn

1. Create an object containing a data frame that is a subset of your name.
2. Create a new percentage variable 'perc', where $\text{prop} * 100$ *hint: mutate()*
3. Create a plot of the popularity of your name over time. Weird trends? Do you need to subset again?
4. Reorder the rows from descending from highest to lowest by perc. What year was most popular for your name? *hint: arrange()*
5. Reorder by year. What were the most popular names in your birth year?


The background image shows a busy city street, likely in London, with a large crowd of people walking. Tall buildings line the street, and a traffic light is visible on the left. The scene is slightly blurred, emphasizing the movement and density of the urban environment.

Crime Dataset


Is there a relationship between crime and temperature? State statistics from 2009.

Five variables: state, abbr, low, murder, tc2009

Describe this relationship
with a single line
What does the line look like?



What does the line look like?




The linear regression algorithm constrains $f(\hat{x})$ to have the form,


$$f(\hat{x}) = \alpha + \beta x + \epsilon$$

e.g., $f(\hat{x})$ will be a straight line in x .

How do we fit the best line?




How do we fit the best line?



```
#read in crime dataset
```


```
crime <- read.csv("crime.csv")
```

```
head(crime)
```



```
mod <- lm(tc2009 ~ low, data = crime)
```

```
qplot(low, tc2009, data=crime)
```



```
mod <- lm(tc2009 ~ low, data=crime)
mod
```

```
## Call:
```

```
## lm(formula = tc2009 ~ low, data = crime)
```

```
## Coefficients:
```

```
## (Intercept)      low
```

```
##  491.135      3.002
```


$$y = \alpha + \beta x + \epsilon$$

α is the expected value of y when x is 0.

β is the expected increase in y associated with a one unit increase in x

coefficients(mod)

(Intercept)	low
491.13	3.002

coefficients(mod)

(Intercept)	low
491.13	3.002
α	β

coefficients(mod)

(Intercept)	low
491.13	3.002
α	β

The best estimate of tc2009 for a state with low = -10 is,

$$491.13 + 3.002 * (-10) = 461.11$$

Extracting info

A common pattern for R models: [store and explore](#)

1. Create model object
2. Run function(s) on model object

summary()
predict()
resid()
plot()

summary(mod)

Call:

```
lm(formula = tc2009 ~ low, data = crime)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-287.96	-88.37	-16.44	62.13	797.60

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	491.135	59.588	8.242	8.16e-11 ***
low	3.002	1.365	2.200	0.0326 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 172.2 on 49 degrees of freedom

Multiple R-squared: 0.08986, Adjusted R-squared: 0.07129

F-statistic: 4.838 on 1 and 49 DF, p-value: 0.03259

summary(mod)

```
Call:
lm(formula = tc2009 ~ low, data = crime)


Residuals:
    Min       1Q   Median       3Q      Max
-287.96  -88.37  -16.44   62.13  797.60

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  491.135     59.588   8.242 8.16e-11 ***
low           3.002       1.365   2.200  0.0326 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 172.2 on 49 degrees of freedom
Multiple R-squared:  0.08986,    Adjusted R-squared:  0.07129
F-statistic: 4.838 on 1 and 49 DF,  p-value: 0.03259
```


Does the model do better at predicting Y than random chance?

```
qplot(low, tc2009, data = crime) +  
geom_smooth(method = "lm")
```





```
qplot(low, tc2009, data = crime) +  
geom_smooth(se = FALSE, method = "lm")
```







`plot(mod)`







`plot(mod)`



plot(mod)



`plot(mod)`




```
par(mfrow = c(2,2))  
plot(mod)
```

Heights data set

Earnings vs. height and demographic characteristics of 1,192 individuals, collected in 1994

Sampled from Gelman and Hill. *Data Analysis using Regression and Multilevel/Hierarchical Models*.
Cambridge Press, 2007.




Your Turn

Read in the “heights.csv” file as an object.

Regress `earn` on `height` and name this linear model object *m1*

Discuss your interpretation of the height coefficient with your neighbor



```
m1 <- lm(earn ~ height, data = heights)
```

```
m1 <- lm(earn ~ height, data = heights)
summary(m1)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-30043 -11422  -3608    6443 173488

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -58611.9     9525.6  -6.153 1.04e-09 ***
height       1221.9       142.1   8.598  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18900 on 1190 degrees of freedom
Multiple R-squared:  0.05849, Adjusted R-squared:  0.0577
F-statistic: 73.93 on 1 and 1190 DF,  p-value: < 2.2e-16
```

```
m1 <- lm(earn ~ height, data = heights)
summary(m1)
```

Residuals:

Min	1Q	Median	3Q	Max
-30043	-11422	-3608	6443	173488

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	-58611.9	9525.6	-6.153	1.04e-09	***
height	1221.9	142.1	8.598	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1					

Residual standard error: 18900 on 1190 degrees of freedom
Multiple R-squared: 0.05849, Adjusted R-squared: 0.0577
F-statistic: 73.93 on 1 and 1190 DF, p-value: < 2.2e-16

```
m1 <- lm(earn ~ height, data = heights)
summary(m1)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-30043 -11422  -3608    6443  173488

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept) -58611.9      9525.6   -6.153  1.04e-09 ***
height       1221.9       142.1    8.598  < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18900 on 1190 degrees of freedom
Multiple R-squared:  0.05849, Adjusted R-squared:  0.0577
F-statistic: 73.93 on 1 and 1190 DF,  p-value: < 2.2e-16
```

```
m1 <- lm(earn ~ height, data = heights)
summary(m1)
```

Each 1" increase in height is associated with a \$1,221.90 increase in earnings

Residuals:

Min	1Q	Median	3Q	Max
-30043	-11422	-3608	6443	173488

The best estimate of earn for someone 68 inches tall is,

-58611.9 + 1221.95 * (68) = 24,477.3

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-58611.9	9525.6	-6.153	1.04e-09 ***
height	1221.9	142.1	8.598	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18900 on 1190 degrees of freedom
Multiple R-squared: 0.05849, Adjusted R-squared: 0.0577
F-statistic: 73.93 on 1 and 1190 DF, p-value: < 2.2e-16

```
m2 <- lm(earn ~ height + sex, data = heights)
summary(m2)
```

```
m2 <- lm(earn ~ height + sex, data = heights)
summary(m2)
```

Call:

```
lm(formula = earn ~ height + sex, data = heights)
```

Residuals:

Min	1Q	Median	3Q	Max
-30018	-11127	-3260	6080	170360

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-5732.9	12665.4	-0.453	0.6509
height	371.7	195.7	1.899	0.0578 .
sexmale	9479.7	1526.0	6.212	7.21e-10 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18610 on 1189 degrees of freedom

Multiple R-squared: 0.08809, Adjusted R-squared: 0.08656

F-statistic: 57.43 on 2 and 1189 DF, p-value: < 2.2e-16

```
m2 <- lm(earn ~ height + sex, data = heights)
summary(m2)
```

Regression with categorical variables creates a “reference” category. In this example the reference or baseline is female and the coefficient (sexmale) is rate of change relative to that baseline.

How can we change this to look at it the other way around?

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-5732.9	12665.4	-0.453	0.6509
height	371.7	195.7	1.899	0.0578
sexmale	9479.7	1526.0	6.212	7.21e-10 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18610 on 1189 degrees of freedom
Multiple R-squared: 0.08809, Adjusted R-squared: 0.08656
F-statistic: 57.43 on 2 and 1189 DF, p-value: < 2.2e-16


```
heights$sex <- factor(heights$sex, levels = c("male", "female"))
m3 <- lm(earn ~ height + sex, data = heights)
summary(m3)
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3746.8	13737.5	0.273	0.7851
height	371.7	195.7	1.899	0.0578 .
sexfemale	-9479.7	1526.0	-6.212	7.21e-10 ***


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18610 on 1189 degrees of freedom

Multiple R-squared: 0.08809, Adjusted R-squared: 0.08656

F-statistic: 57.43 on 2 and 1189 DF, p-value: < 2.2e-16

```
qplot(height, earn, color= sex, data = heights) + geom_smooth(method = "lm", se = FALSE)
```



Resources

Cran task views

<http://cran.r-project.org/web/views/>

Categorization of R packages for specific applications.
For example, the "Finance" task view displays packages
which are applicable in Finance.

CRAN Task Views	
Bayesian	Bayesian Inference
ChemPhys	Chemometrics and Computational Physics
ClinicalTrials	Clinical Trial Design, Monitoring, and Analysis
Cluster	Cluster Analysis & Finite Mixture Models
DifferentialEquations	Differential Equations
Distributions	Probability Distributions
Econometrics	Computational Econometrics
Environmetrics	Analysis of Ecological and Environmental Data
ExperimentalDesign	Design of Experiments (DoE) & Analysis of Experimental Data
Finance	Empirical Finance
Genetics	Statistical Genetics

R Bloggers

<http://www.r-bloggers.com/>

Aggregation of R blogs. This is a great resource to read about new things in R and what the R community is doing. You'll often find new tips and tricks to enhance your R skills.

R-bloggers

R news and tutorials contributed by (452) R bloggers

[Home](#)[About](#)[add your blog!](#)[Contact us](#)[RSS](#)[R jobs](#)

WELCOME!

Here you will find daily news and tutorials about R, contributed by over 450 bloggers. You can subscribe for e-mail updates:

11218 readers
BY FEEDBURNER

And get updates to your Facebook:

High Dimensional Biological Data Analysis and Visualization

February 22, 2014

By dgrapov



High dimensional biological data shares many qualities with other forms of data. Typically it is wide (samples << variables), complicated by experiential design and made up of complex

TOP 3 POSTS FROM THE PAST 2 DAYS

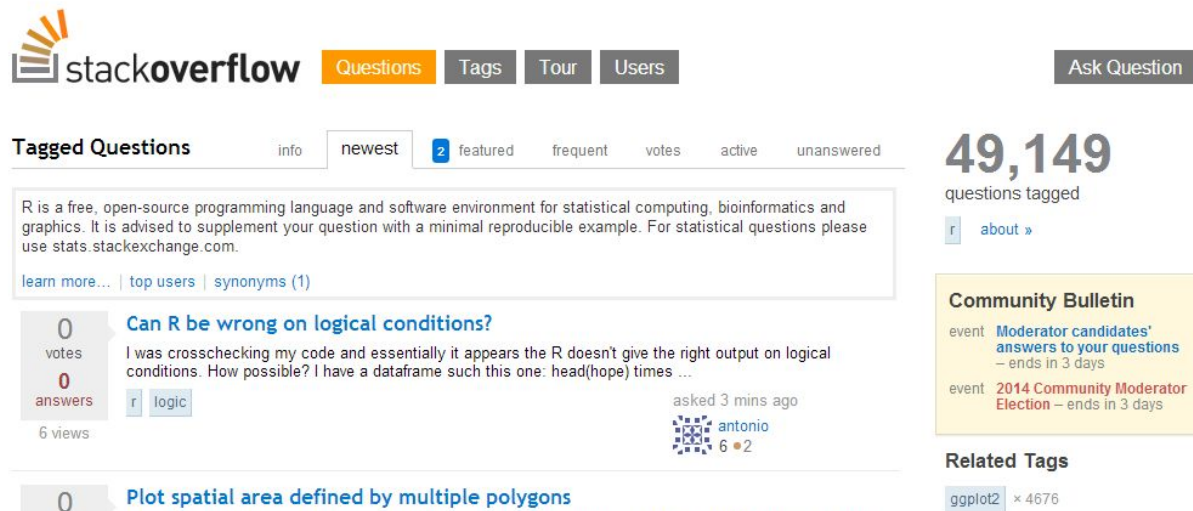
The gap between data mining and predictive models
No need for SPSS – beautiful output in R #rstats
Because it's Friday: US Dialects

TOP 9 ARTICLES OF THE WEEK

Stack Overflow

<http://stackoverflow.com/questions/tagged/r>

The R tag on Stack Overflow is becoming an increasingly important resource for seeking answers to R related questions. You can search the R tag in general, or refine your search to another tag such as ggplot2 or sweave.



The screenshot displays the Stack Overflow interface for the R tag. At the top, the Stack Overflow logo is on the left, and navigation links for Questions, Tags, Tour, and Users are in the center. An 'Ask Question' button is on the right. Below the navigation bar, the 'Tagged Questions' section is active, with tabs for info, newest, featured (2), frequent, votes, active, and unanswered. A descriptive box for the R tag explains it's a free, open-source programming language for statistical computing and bioinformatics, advising users to provide a minimal reproducible example and use stats.stackexchange.com. Below this, a question titled 'Can R be wrong on logical conditions?' is shown, asking about logical conditions in R. It has 0 votes, 0 answers, and 6 views. The question was asked 3 minutes ago by user 'antonio'. To the right, a 'Community Bulletin' section lists two events: 'Moderator candidates' answers to your questions' (ends in 3 days) and '2014 Community Moderator Election' (ends in 3 days). At the bottom, a 'Related Tags' section shows 'ggplot2' with 4676 questions.

stackoverflow Questions Tags Tour Users Ask Question

Tagged Questions info newest 2 featured frequent votes active unanswered

R is a free, open-source programming language and software environment for statistical computing, bioinformatics and graphics. It is advised to supplement your question with a minimal reproducible example. For statistical questions please use stats.stackexchange.com.

learn more... | top users | synonyms (1)

0 votes
0 answers
6 views

Can R be wrong on logical conditions?

I was crosschecking my code and essentially it appears the R doesn't give the right output on logical conditions. How possible? I have a dataframe such this one: head(hope) times ...

asked 3 mins ago
antonio
6 • 2

Community Bulletin

event **Moderator candidates' answers to your questions**
– ends in 3 days

event **2014 Community Moderator Election**
– ends in 3 days

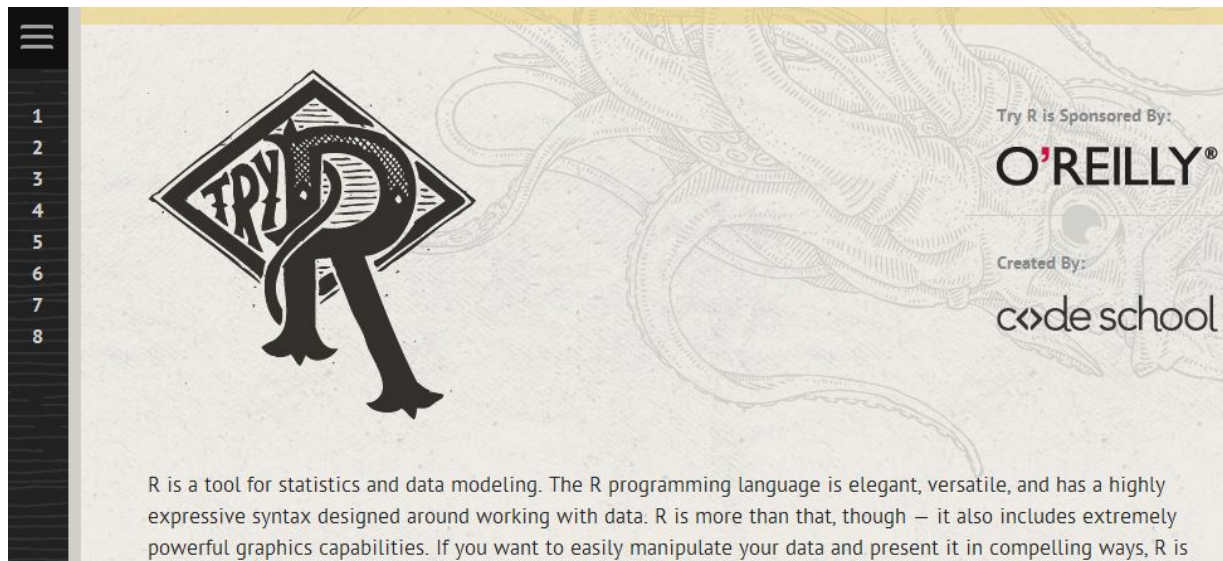
Related Tags

ggplot2 × 4676

try R

<http://tryr.codeschool.com/>

A game-ified R lesson that provides an embedded console and quick feedback, a useful way to practice the basics of R.



R Documentation

<http://www.rdocumentation.org/>

A searchable, enhanced version of R's documentation pages. Includes package download stats and user comments. You can run and manipulate example code right in the webpage.

The screenshot shows the R Documentation website. On the left is a sidebar with a search bar and a list of domains. The main content area has a search bar, a description of the tool, and search input fields. On the right is a 'Top Ranked Packages' section with a table of popular packages.

DOMAINS

- Bayesian
- ChemPhys
- ClinicalTrials
- Cluster
- DifferentialEquations
- Econometrics
- Finance
- Genetics
- Graphics
- HighPerformanceComputing
- MachineLearning
- MedicalImaging
- MetaAnalysis
- Multivariate
- NaturalLanguageProcessing

R Documentation

Search the R documentation of **5340** R packages and **111209** R functions:

Rdocumentation is a tool that helps you easily find and browse the documentation of all current and some past packages on CRAN. Click on the search bar at the top left for instant search or fill out the forms below for advanced search!

All Fields

Package Name

Function Name

Title

Rdocumentation package

Top Ranked Packages



Week | Month | All time

#	Package	#d
1 -	digest	42844
2 -	ggplot2	39696
3 -	RColorBrewer	39185
4 -	plyr	38241
5 -	stringr	38068
6 -	colorspace	34851
7 -	survival	33024

Quick-R

<http://www.statmethods.net/>

Excellent, “quick” reference to common R operations.



ggplot2 doc

<http://docs.ggplot2.org/current/>

ggplot2 documentation organized by components of “grammar of graphics”

ggplot2 0.9.3.1 [Index](#)

Help topics

Geoms

Geoms, short for geometric objects, describe the type of plot you will produce.

- `geom_abline`
Line specified by slope and intercept.
- `geom_area`
Area plot.
- `geom_bar`
Bars, rectangles with bases on x-axis
- `geom_bin2d`
Add heatmap of 2d bin counts.
- `geom_blank`




Dependencies

- **Depends:** stats, methods
- **Imports:** plyr, digest, grid, gtable, reshape2, scales, proto, MASS
- **Suggests:** quantreg, Hmisc, mapproj, maps, hexbin, maptools, multcomp, nlme, testthat
- **Extends:**

Library Research Guide

<http://researchguides.library.tufts.edu/data/r>

Library links to O'Reilly ebooks. Free to the Tufts community!




[Qualitative Comparative Analysis with R](#) - Alrik Thiem; Adrian Duda

ISBN: 9781461445838

Publication Date: 2012-08-30

Social science theory often builds on sets and their relations. Correlation-based methods of scientific enquiry, however, use linear algebra and are unsuited to analyzing set relations. The development of Qualitative Comparative Analysis (QCA) by Charles Ragin has given social scientists a formal tool for identifying set-theoretic connections based on Boolean algebra. As a result, interest in this method has markedly risen among social scientists in recent years. This book offers the first complete introduction on how to perform QCA in the R software environment for statistical computing and graphics with the QCA package.




[R Cookbook](#) - Paul Teetor

ISBN: 9780596809157

Publication Date: 2011-03-22

With more than 200 practical recipes, this book helps you perform data analysis with R quickly and efficiently. The R language provides everything you need to do statistical work, but its structure can be difficult to master.



[R in a Nutshell](#) - Joseph Adler

ISBN: 144931208X

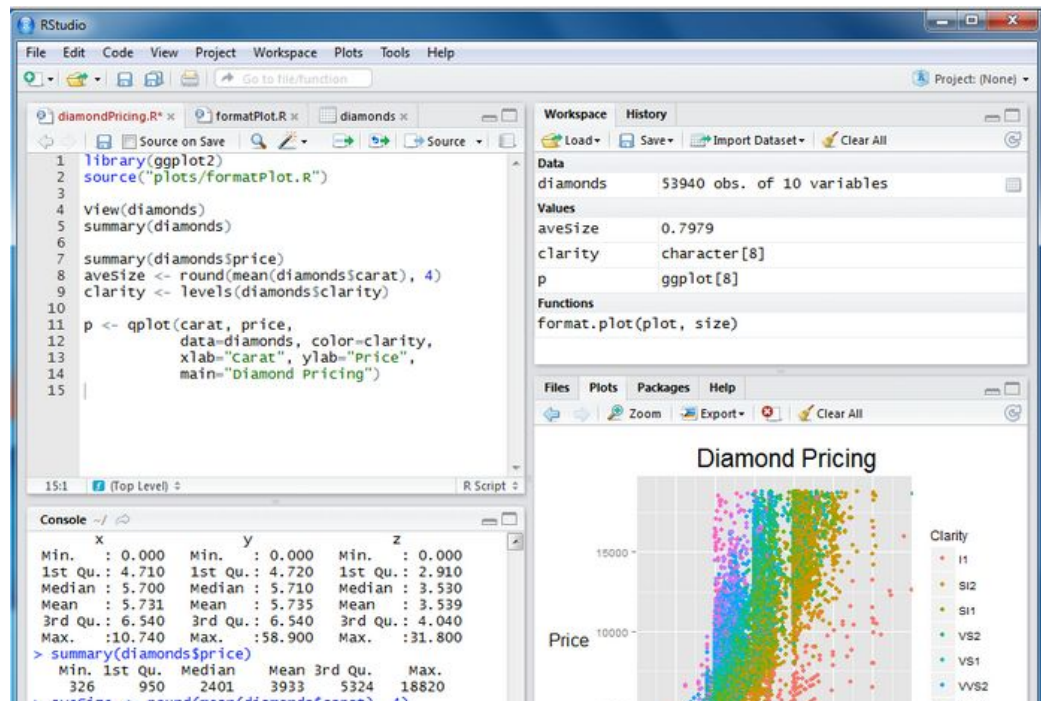
Publication Date: 2012-10-16

If you're considering R for statistical computing and data visualization, this book provides a quick and practical guide to just about everything you can do with the open source R language and software environment. You'll learn how to write R functions and use R packages to help you prepare, visualize, and analyze data.

RStudio

<https://www.rstudio.com/>

A powerful and convenient user environment for R. Free and open source. Highly, highly recommended.



To be successful with R

- Like learning any programming language, take your time and try to run the code in your head before you run it on your machine. Try to predict what will happen.
- Be patient.
- Think of a fun project that you actually would like to do and do it in R.
- Ask your friendly librarian about useful resources.
- Understand that unlike excel, there are many paths to the same solution in R. You need not learn them all but to troubleshoot effectively and ask for help it is worthwhile to understand how others might work with R (ie., subsetting)
- If you master the techniques and concepts in this workshop you've mastered 80% of R. The rest is identifying specific packages/methodologies that are relevant to your [domain](#).

thanks!

contact: datahelp@tufts.edu